

In [1]: # importing modules

```
import os
import pandas as pd
import numpy as np
from glob import glob
import tmdbsimple as tmdb
import requests
```

executed in 299ms, finished 13:39:58 2021-05-01

In [2]: csv_list = glob("./dataFiles/*.csv")

```
csv_list
```

executed in 12ms, finished 13:39:58 2021-05-01

Out[2]: ['./dataFiles\bom.movie_gross.csv',

```
'./dataFiles\expanded_tmdb.csv',
'./dataFiles\movie_collection.csv',
'./dataFiles\name.basics.csv',
'./dataFiles\title.akas.csv',
'./dataFiles\title.basics.csv',
'./dataFiles\title.crew.csv',
'./dataFiles\title.principals.csv',
'./dataFiles\title.ratings.csv',
'./dataFiles\tmdb.movies.csv',
'./dataFiles\tn.movie_budgets.csv']
```

In [3]: cleaned_filenames = [filename.split('\\')[1].replace('.csv', '').replace('.', '_')
for filename in csv_list]

executed in 18ms, finished 13:39:59 2021-05-01

In [4]: orig_dfs = {}

```
for idx, file in enumerate(csv_list):
    orig_dfs.update({cleaned_filenames[idx]: pd.read_csv(file)})
```

executed in 2.02s, finished 13:40:01 2021-05-01

In [5]: orig_dfs['bom_movie_gross'].head()

executed in 14ms, finished 13:40:01 2021-05-01

Out[5]:

		title	studio	domestic_gross	foreign_gross	year
0		Toy Story 3	BV	415000000.0	652000000	2010
1	Alice in Wonderland (2010)		BV	334200000.0	691300000	2010
2	Harry Potter and the Deathly Hallows Part 1		WB	296000000.0	664300000	2010
3	Inception		WB	292600000.0	535700000	2010
4	Shrek Forever After		P/DW	238700000.0	513900000	2010

Now that we have all the files loaded into dataframes in an easily accessible structure, lets investigate them all to see what information we have available and make connections between the data included in each file. Besides, we can see what modifications we may need to make to data

types, to account for missing values, etc. Going forward we will get a lot of information to keep track of, so I will make a separate excel doc for use as a reference / table schema to relate these dataframes in the future.

```
In [6]: for item in orig_dfs.items():
    print("=====New DF====")
    print(f'Title: {item[0]}')
    print(item[1].info())
    print("=====End DF====")
```

executed in 320ms, finished 13:40:02 2021-05-01

```
=====New DF=====
Title: bom_movie_gross
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3387 entries, 0 to 3386
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   title        3387 non-null   object  
 1   studio       3382 non-null   object  
 2   domestic_gross 3359 non-null   float64 
 3   foreign_gross 2037 non-null   object  
 4   year         3387 non-null   int64  
dtypes: float64(1), int64(1), object(3)
memory usage: 132.4+ KB
None
=====End DF=====
=====New DF=====
Title: expanded_tmdb
<class 'pandas.core.frame.DataFrame'>
```

So far, we can notice that, based on column names, there are some potential connections between the different dataframes. Additionally, there are columns that are hard to interpret from their labels, and there are some missing a large majority of values. Lastly, we can note that there are columns with data types that might be undesirable, such as gross columns stored as strings, and date columns stored as objs or ints/floats. First, to get a better understanding of what is in each column, lets look at the head of each dataframe (since `print(df.head())` doesn't look as nice in the notebook, create cell for each one instead).

```
In [7]: orig_dfs['bom_movie_gross'].head()
```

executed in 15ms, finished 13:40:02 2021-05-01

Out[7]:

		title	studio	domestic_gross	foreign_gross	year
0		Toy Story 3	BV	415000000.0	652000000	2010
1		Alice in Wonderland (2010)	BV	334200000.0	691300000	2010
2		Harry Potter and the Deathly Hallows Part 1	WB	296000000.0	664300000	2010
3		Inception	WB	292600000.0	535700000	2010
4		Shrek Forever After	P/DW	238700000.0	513900000	2010

In [8]: `orig_dfs['name_basics'].head()`

executed in 14ms, finished 13:40:02 2021-05-01

Out[8]:

	nconst	primary_name	birth_year	death_year	primary_profession
0	nm0061671	Mary Ellen Bauder	NaN	NaN	miscellaneous,production_manager,producer
1	nm0061865	Joseph Bauer	NaN	NaN	composer,music_department,sound_department
2	nm0062070	Bruce Baum	NaN	NaN	miscellaneous,actor,writer
3	nm0062195	Axel Baumann	NaN	NaN	camera_department,cinematographer,art_department
4	nm0062798	Pete Baxter	NaN	NaN	production_designer,art_department,set_decoration

In [9]: `orig_dfs['title_akas'].head()`

executed in 15ms, finished 13:40:02 2021-05-01

Out[9]:

	title_id	ordering	title	region	language	types	attributes	is_original_title
0	tt0369610	10	Джурасик свят	BG	bg	NaN	NaN	0.0
1	tt0369610	11	Jurashikku warudo	JP	NaN	imdbDisplay	NaN	0.0
2	tt0369610	12	Jurassic World: O Mundo dos Dinossauros	BR	NaN	imdbDisplay	NaN	0.0
3	tt0369610	13	O Mundo dos Dinossauros	BR	NaN	NaN	short title	0.0
4	tt0369610	14	Jurassic World	FR	NaN	imdbDisplay	NaN	0.0

In [10]: `orig_dfs['title_basics'].head()`

executed in 15ms, finished 13:40:03 2021-05-01

Out[10]:

	tconst	primary_title	original_title	start_year	runtime_minutes	genres
0	tt0063540	Sunghursh	Sunghursh	2013	175.0	Action,Crime,Drama
1	tt0066787	One Day Before the Rainy Season	Ashad Ka Ek Din	2019	114.0	Biography,Drama
2	tt0069049	The Other Side of the Wind	The Other Side of the Wind	2018	122.0	Drama
3	tt0069204	Sabse Bada Sukh	Sabse Bada Sukh	2018	NaN	Comedy,Drama
4	tt0100275	The Wandering Soap Opera	La Telenovela Errante	2017	80.0	Comedy,Drama,Fantasy

In [11]: `orig_dfs['title_crew'].head()`

executed in 14ms, finished 13:40:03 2021-05-01

Out[11]:

	tconst	directors	writers
0	tt0285252	nm0899854	nm0899854
1	tt0438973	NaN	nm0175726,nm1802864
2	tt0462036	nm1940585	nm1940585
3	tt0835418	nm0151540	nm0310087,nm0841532
4	tt0878654	nm0089502,nm2291498,nm2292011	nm0284943

In [12]: `orig_dfs['title_principals'].head()`

executed in 14ms, finished 13:40:03 2021-05-01

Out[12]:

	tconst	ordering	nconst	category	job	characters
0	tt0111414	1	nm0246005	actor	NaN	["The Man"]
1	tt0111414	2	nm0398271	director	NaN	NaN
2	tt0111414	3	nm3739909	producer	producer	NaN
3	tt0323808	10	nm0059247	editor	NaN	NaN
4	tt0323808	1	nm3579312	actress	NaN	["Beth Boothby"]

In [13]: `orig_dfs['title_ratings'].head()`

executed in 14ms, finished 13:40:03 2021-05-01

Out[13]:

	tconst	averagerating	numvotes
0	tt10356526	8.3	31
1	tt10384606	8.9	559
2	tt1042974	6.4	20
3	tt1043726	4.2	50352
4	tt1060240	6.5	21

In [14]: `orig_dfs['tmdb_movies'].head()`

executed in 15ms, finished 13:40:04 2021-05-01

Out[14]:

	Unnamed: 0	genre_ids	id	original_language	original_title	popularity	release_date	title
0	0	[12, 14, 10751]	12444	en	Harry Potter and the Deathly Hallows: Part 1	33.533	2010-11-19	Harry Potter and the Deathly Hallows: Part 1
1	1	[14, 12, 16, 10751]	10191	en	How to Train Your Dragon	28.734	2010-03-26	How to Train Your Dragon
2	2	[12, 28, 878]	10138	en	Iron Man 2	28.515	2010-05-07	Iron Man 2
3	3	[16, 35, 10751]	862	en	Toy Story	28.005	1995-11-22	Toy Story
4	4	[28, 878, 12]	27205	en	Inception	27.920	2010-07-16	Inception

In [15]: `orig_dfs['tn_movie_budgets'].head(10)`

executed in 15ms, finished 13:40:04 2021-05-01

Out[15]:

	id	release_date	movie	production_budget	domestic_gross	worldwide_gross
0	1	Dec 18, 2009	Avatar	\$425,000,000	\$760,507,625	\$2,776,345,279
1	2	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	\$410,600,000	\$241,063,875	\$1,045,663,875
2	3	Jun 7, 2019	Dark Phoenix	\$350,000,000	\$42,762,350	\$149,762,350
3	4	May 1, 2015	Avengers: Age of Ultron	\$330,600,000	\$459,005,868	\$1,403,013,963
4	5	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	\$317,000,000	\$620,181,382	\$1,316,721,747
5	6	Dec 18, 2015	Star Wars Ep. VII: The Force Awakens	\$306,000,000	\$936,662,225	\$2,053,311,220
6	7	Apr 27, 2018	Avengers: Infinity War	\$300,000,000	\$678,815,482	\$2,048,134,200
7	8	May 24, 2007	Pirates of the Caribbean: At World's End	\$300,000,000	\$309,420,425	\$963,420,425
8	9	Nov 17, 2017	Justice League	\$300,000,000	\$229,024,295	\$655,945,209
9	10	Nov 6, 2015	Spectre	\$300,000,000	\$200,074,175	\$879,620,923

Seeing these, there are some obvious joins / connections to make on things we can be nearly confident in referencing the same thing, such as on identifiers tconst and nconst. There are some columns we are somewhat confident match but will have to be careful about, such as the different

movie title columns, and some columns we are not sure yet referencing the same movie (or have equivalent values if they do) but potentially could be, such as the ratings and quoted domestic and foreign gross values. Lastly, there are some extra indexing columns we can drop, such as unnamed and id in the last two tables. The ordering columns as well perhaps, as they seem to index rows for multiple records for the same title. Let's get rid of cluttering rows we know aren't useful first. Then we can look at combining dataframes / eliminating other columns.

```
In [16]: orig_dfs['tmdb_movies'].drop(columns = ['Unnamed: 0'], inplace = True)
orig_dfs['tn_movie_budgets'].drop(columns = ['id'], inplace = True)
orig_dfs['title_principals'].drop(columns = ['ordering'], inplace = True)
orig_dfs['title_akas'].drop(columns = ['ordering'], inplace = True)
```

executed in 46ms, finished 13:40:04 2021-05-01

Now, let's look at other problem columns and fix them up. Starting with bom_movie_gross and moving down the line, we see that in that dataframe there is a studio column that could be more interpretable, since they are currently abbreviations. Let's first see what all the unique abbreviations are.

In [17]: `orig_dfs['bom_movie_gross'].studio.unique()`

executed in 14ms, finished 13:40:04 2021-05-01

Out[17]: `array(['BV', 'WB', 'P/DW', 'Sum.', 'Par.', 'Uni.', 'Fox', 'Wein.', 'Sony', 'FoxS', 'SGem', 'WB (NL)', 'LGF', 'MBox', 'CL', 'W/Dim.', 'CBS', 'Focus', 'MGM', 'Over.', 'Mira.', 'IFC', 'CJ', 'NM', 'SPC', 'ParV', 'Gold.', 'JS', 'RAtt.', 'Magn.', 'Free', '3D', 'UTV', 'Rela.', 'Zeit.', 'Anch.', 'PDA', 'Lorb.', 'App.', 'Drft.', 'Osci.', 'IW', 'Rog.', nan, 'Eros', 'Relbig.', 'Viv.', 'Hann.', 'Strand', 'NGE', 'Scre.', 'Kino', 'Abr.', 'CZ', 'ATO', 'First', 'GK', 'FInd.', 'NFC', 'TFC', 'Pala.', 'Imag.', 'NAV', 'Arth.', 'CLS', 'Mont.', 'Olive', 'CGld', 'FOAK', 'IVP', 'Yash', 'ICir', 'FM', 'Vita.', 'WOW', 'Truly', 'Indic.', 'FD', 'Vari.', 'TriS', 'ORF', 'IM', 'Elev.', 'Cohen', 'NeoC', 'Jan.', 'MNE', 'Trib.', 'Rocket', 'OMNI/FSR', 'KKM', 'Argo.', 'SMod', 'Libre', 'FRun', 'WHE', 'P4', 'KC', 'SD', 'AM', 'MPFT', 'Icar.', 'AGF', 'A23', 'Da.', 'NYer', 'Rialto', 'DF', 'KL', 'ALP', 'LG/S', 'WGUSA', 'MPI', 'RTWC', 'FIP', 'RF', 'ArcEnt', 'PalUni', 'EpicPics', 'EOne', 'LD', 'AF', 'TFA', 'Myr.', 'BM&DH', 'SEG', 'Palt', 'Outs', 'OutF', 'BSM', 'WAMCR', 'PM&E', 'A24', 'Cdgm.', 'Distrib.', 'Imax', 'PH', 'HTR', 'ELS', 'PI', 'E1', 'TVC', 'FEF', 'EXCL', 'MSF', 'P/108', 'FCW', 'XL', 'Shout!', 'SV', 'CE', 'VPD', 'KE', 'Saban', 'CF&SR', 'Triu', 'DR', 'Crnths', 'Ampl.', 'CP', 'Proud', 'BGP', 'Abk.', 'DLA', 'B360', 'BWP', 'SEA', 'RME', 'KS', 'VE', 'LGP', 'EC', 'FUN', 'STX', 'AR', 'BG', 'PFR', 'BST', 'BH Tilt', 'BSC', 'U/P', 'UHE', 'CLF', 'FR', 'AaF', 'Orch.', 'Alc', 'PBS', 'SHO', 'Grav.', 'Gathr', 'Asp.', 'ADC', 'Rel.', 'SM', 'AZ', 'UEP', 'ITL', 'TA', 'MR', 'BBC', 'CFilms', 'Part.', 'FOR', 'TAFC', 'JBG', 'PNT', 'CineGalaxy', 'Fathom', 'Zee', 'Men.', 'YFG', 'Gaatri', 'Mon', 'Ghop', 'Cleopatra', 'Dreamwest', 'SDS', 'Linn', 'Electric', 'Jampa', 'HC', 'GrtIndia', 'Neon', 'ENTMP', 'Good Deed', 'ParC', 'Aviron', 'Annapurna', 'Amazon', 'Affirm', 'MOM', 'Orion', 'CFI', 'UTMW', 'Crimson', 'CAVU', 'EF', 'Arrow', 'Hiber', 'Studio 8', 'Global Road', 'Trafalgar', 'Greenwich', 'Spanglish', 'Blue Fox', 'RLJ', 'Swen', 'PackYourBag', 'Gaum.', 'Grindstone', 'Conglomerate', 'MUBI', 'Darin Southa', 'Super', 'CARUSEL', 'PDF', 'Synergetic'], dtype=object)`

Okay, well, there's a lot. Most or some of them probably do not exist anymore.. maybe we will wait on this.

A useful discovery from the The Movie Database API page for turning genre_id to interpretable strings in that table: MOVIES Action 28 Adventure 12 Animation 16 Comedy 35 Crime 80 Documentary 99 Drama 18 Family 10751 Fantasy 14 History 36 Horror 27 Music 10402 Mystery 9648 Romance 10749 Science Fiction 878 TV Movie 10770 Thriller 53 War 10752 Western 37

We could use this to replace the genre ID's with text.

Now, we may be able to condense the number of tables while leaving appropriate opportunity for joins. We just need to start checking where records are truly pointing toward the same information, then bringing information together. For example, we can imagine for a unique movie, the following information would be useful to collect for generating visualizations: tconst(table primary key), primary_title (for display purposes), genre(s), release_year, run_time, production_studio,

production_budget, domestic_gross, domestic_ROI, worldwide_gross, worldwide_ROI. We can collect the rest of the useful data in two other tables probably. Brainstorming, I would be interested in one table having columns tconst, rating, num_votes and popularity, and a second table having tconst, and then columns for each crew job type (producer, director, writer, actor, etc.) containing lists of the crew of that type for that film. Let's look at the sizes of each dataframe to see if there are ones we can join to / build off of while retaining the maximum amount of information.

```
In [18]: for item in orig_dfs.items():
    print("=====New DF====")
    print(f'Title: {item[0]}')
    print(item[1].shape)
    print("=====End DF====")
```

executed in 14ms, finished 13:40:05 2021-05-01

```
=====New DF=====
Title: bom_movie_gross
(3387, 5)
=====End DF=====
=====New DF=====
Title: expanded_tmdb
(26517, 15)
=====End DF=====
=====New DF=====
Title: movie_collection
(2373, 12)
=====End DF=====
=====New DF=====
Title: name_basics
(606648, 6)
=====End DF=====
=====New DF=====
Title: title_akas
(331703, 7)
=====End DF=====
=====New DF=====
Title: title_basics
(146144, 6)
=====End DF=====
=====New DF=====
Title: title_crew
(146144, 3)
=====End DF=====
=====New DF=====
Title: title_principals
(1028186, 5)
=====End DF=====
=====New DF=====
Title: title_ratings
(73856, 3)
=====End DF=====
=====New DF=====
Title: tmdb_movies
(26517, 9)
=====End DF=====
=====New DF=====
Title: tn_movie_budgets
(5782, 5)
=====End DF=====
```

So we are seeing that we have data for many more movies than we have budget information for. Since we are trying to produce profitable movies, being able to determine indicators driving profit is necessary, unless we can become very confident that something is an effective proxy for profit such as rating or popularity. It is possible that between the tables with budgetary information, i.e.

'bom_movie_gross' and 'tn_movie_budgets' we can cover more movies than either alone, and hopefully fill out some of the other information with the other tables we have. However, we only have the budget information for those from the 'tn_movie_budgets' table, which means we can only calculate ROI's for movies from that table, unless we can find the information elsewhere OR are happy filling in values for NULLs. We can look into that but it is hard to imagine that budget can easily be predicted and its assigned value has a huge impact on ROI.

Upon doing some research, I have determined by looking at lists on the The Numbers and The Movie Database websites that the domestic, foreign, and worldwide gross numbers from each site are close enough that they can be considered the same. Additionally, the 'worldwide_gross' from the 'tn_movie_budgets' table is the same as the sum of the 'domestic' and 'foreign_gross' columns of the 'bom_movie_gross' table, so we can make an aggregate column in the Box Office Mojo table and also produce a 'foreign_gross' column for the The Numbers table. Down the line, it may be possible to scrape The Numbers or Box Office Mojo for budgetary information for more movies.

Additionally, there is a way to query TMDB using an IMDB tconst through a python module called tmdbsimple:

```
import tmdbsimple as tmdb
tmdb.API_KEY = "response ="
tmdb.Find('tt0266543').info(external_source='imdb_id') response

{'movie_results': [{"id": 12, "video": False, "vote_count": 13478, "vote_average": 7.8, "title": "Finding Nemo", "release_date": "2003-05-30", "original_language": "en", "original_title": "Finding Nemo", "genre_ids": [16, 10751], "backdrop_path": "/dFYguAfeVt19qAbzJ5mArn7DEJw.jpg", "adult": False, "overview": "Nemo, an adventurous young clownfish, is unexpectedly taken from his Great Barrier Reef home to a dentist's office aquarium. It's up to his worrisome father Marlin and a friendly but forgetful fish Dory to bring Nemo home -- meeting vegetarian sharks, surfer dude turtles, hypnotic jellyfish, hungry seagulls, and more along the way.", "poster_path": "/xVNSgrsvpcAHPnyKf2phYxyppNZ.jpg", "popularity": 31.27}], "person_results": [], "tv_results": [], "tv_episode_results": [], "tv_season_results": []}
```

Or, simply using a get request you can get the above info and more, seen at <https://developers.themoviedb.org/3/movies/get-movie-details> (<https://developers.themoviedb.org/3/movies/get-movie-details>).

It may be worth looking into that package or otherwise requesting from the web API to get a complete dataset from one source.

Anyway, let's begin by joining the two tables we have with budgetary information and see what we get.

```
In [19]: budget_merge = orig_dfs['tn_movie_budgets'].merge(orig_dfs['bom_movie_gross'],
                                                       how='left',
                                                       left_on='movie',
                                                       right_on='title')

budget_merge.head(20)
```

executed in 30ms, finished 13:40:05 2021-05-01

Out[19]:

	release_date	movie	production_budget	domestic_gross_x	worldwide_gross	title
0	Dec 18, 2009	Avatar	\$425,000,000	\$760,507,625	\$2,776,345,279	NaN
1	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	\$410,600,000	\$241,063,875	\$1,045,663,875	Pirates of the Caribbean: On Stranger Tides
2	Jun 7, 2019	Dark Phoenix	\$350,000,000	\$42,762,350	\$149,762,350	NaN
3	May 1, 2015	Avengers: Age of Ultron	\$330,600,000	\$459,005,868	\$1,403,013,963	Avengers: Age of Ultron
4	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	\$317,000,000	\$620,181,382	\$1,316,721,747	NaN
5	Dec 18, 2015	Star Wars Ep. VII: The Force Awakens	\$306,000,000	\$936,662,225	\$2,053,311,220	NaN
6	Apr 27, 2018	Avengers: Infinity War	\$300,000,000	\$678,815,482	\$2,048,134,200	Avengers: Infinity War
7	May 24, 2007	Pirates of the Caribbean: At Worldâ„¢s End	\$300,000,000	\$309,420,425	\$963,420,425	NaN
8	Nov 17, 2017	Justice League	\$300,000,000	\$229,024,295	\$655,945,209	Justice League
9	Nov 6, 2015	Spectre	\$300,000,000	\$200,074,175	\$879,620,923	Spectre
10	Jul 20, 2012	The Dark Knight Rises	\$275,000,000	\$448,139,099	\$1,084,439,099	The Dark Knight Rises
11	May 25, 2018	Solo: A Star Wars Story	\$275,000,000	\$213,767,512	\$393,151,347	Solo: A Star Wars Story
12	Jul 2, 2013	The Lone Ranger	\$275,000,000	\$89,302,115	\$260,002,115	The Lone Ranger
13	Mar 9, 2012	John Carter	\$275,000,000	\$73,058,679	\$282,778,100	John Carter
14	Nov 24, 2010	Tangled	\$260,000,000	\$200,821,936	\$586,477,240	Tangled
15	May 4, 2007	Spider-Man 3	\$258,000,000	\$336,530,303	\$894,860,230	NaN

	release_date	movie	production_budget	domestic_gross_x	worldwide_gross	title
16	May 6, 2016	Captain America: Civil War	\$250,000,000	\$408,084,349	\$1,140,069,413	Captain America: Civil War
17	Mar 25, 2016	Batman v Superman: Dawn of Justice	\$250,000,000	\$330,360,194	\$867,500,281	Batman v Superman: Dawn of Justice
18	Dec 14, 2012	The Hobbit: An Unexpected Journey	\$250,000,000	\$303,003,568	\$1,017,003,568	The Hobbit: An Unexpected Journey
19	Jul 15, 2009	Harry Potter and the Half-Blood Prince	\$250,000,000	\$302,089,278	\$935,213,767	NaN

In [20]: budget_merge.shape

executed in 15ms, finished 13:40:05 2021-05-01

Out[20]: (5782, 10)

To be honest, after looking at this stuff, it is very tempting to take the tmdb movies dataframe and use those IDs or otherwise generate a completely new and much larger dataset from TMDB. It can result in having budget information for way more movies. However, I notice that their budget estimates are quite low compared to The Numbers estimates, and seeing as both are really guesses and are user maintained / reported it is hard to know which to be more confident about. So long as we are consistent, it should be fine. I think, as opposed to replacing the entire dataset, we can use their API to get a few columns of additional information for the tmdb_movies dataframe, i.e. their imdb ID number (eg. tt#####) for connecting reliably with other tables, maybe the production studio, and the listed budget and revenue for each film, since that will solve the problem of having to look at an intersection of the information collected from Box Office Mojo and The Numbers.

We can do this very easily using the tmdbsimple python module I mentioned above. First, we had to create an account and request an API key for TMDB. For safekeeping, here is my API key for TMDB: 0c49f68468be72ef8df12e077b3ab923

Below is an example of how to get the budget for a movie (Harry Potter: DH1) we have in our tmdb dataframe:

```
tmdb.API_KEY = '0c49f68468be72ef8df12e077b3ab923'
test_movie = tmdb.Movies(12444)
response = test_movie.info()
```

executed in 62ms, finished 13:40:05 2021-05-01

```
In [22]: print(test_movie.runtime)
print(test_movie.budget)
print(test_movie.production_companies[0]['name'])
print(test_movie.revenue)
print(test_movie.vote_average)
print(test_movie.vote_count)
print(test_movie.imdb_id)
```

executed in 15ms, finished 13:40:06 2021-05-01

```
146
250000000
Warner Bros. Pictures
954305868
7.8
14487
tt0926084
```

Let's use these values to verify equivalency between this information and some in the other dataframes we already have.

```
In [23]: orig_dfs['title_basics'].loc[orig_dfs['title_basics']['tconst'] == 'tt0926084']
```

executed in 14ms, finished 13:40:06 2021-05-01

Out[23]:

	tconst	primary_title	original_title	start_year	runtime_minutes	genres
457	tt0926084	Harry Potter and the Deathly Hallows: Part 1	Harry Potter and the Deathly Hallows: Part 1	2010	146.0	Adventure,Fantasy,Mystery

```
In [24]: orig_dfs['title_ratings'].loc[orig_dfs['title_ratings']['tconst'] == 'tt0926084']
```

executed in 12ms, finished 13:40:06 2021-05-01

Out[24]:

	tconst	averagerating	numvotes
65053	tt0926084	7.7	425530

```
In [25]: orig_dfs['tn_movie_budgets'].loc[orig_dfs['tn_movie_budgets']['movie'] == 'Harry Po
```

executed in 15ms, finished 13:40:07 2021-05-01

Out[25]:

	release_date	movie	production_budget	domestic_gross	worldwide_gross
--	--------------	-------	-------------------	----------------	-----------------

```
In [26]: orig_dfs['bom_movie_gross'].loc[orig_dfs['bom_movie_gross']['title'] == 'Harry Po
```

executed in 15ms, finished 13:40:07 2021-05-01

Out[26]:

	title	studio	domestic_gross	foreign_gross	year
2	Harry Potter and the Deathly Hallows Part 1	WB	296000000.0	664300000	2010

Okay, we are seeing conformance in some and not in others, by looking at these records and also

manually comparing some values on the actual websites. Notice that grabbing the imdb id allows us to easily connect to other tables as it is indeed the same as tconst, as expected. Likewise, while not exactly equal, TMDB's revenue info seems to be the same as worldwide_gross in tn_movie_budgets (and the same as domestic plus foreign gross in bom_movie_gross, were the record to exist). The production studio appears to be the same, however TMDB will return a list of companies that the name needs to be grabbed from (it may work to just grab the first company in the list if there is more than one as in the example above). Finally, we previously recognize that estimated budgets from TMDB are lower, sometimes much lower than the estimates from The Numbers. However, it is probably better to consistently use values from TMDB than to fill in missing values when joining the other two tables with financial information and limiting the size of available data significantly. On the other hand, we note that, compared to the genres from IMDB, the ones from TMDB aren't necessarily the same, though they may include common values. It may be worth it to compare the two columns and keep like genres as it is most likely those are the primary genre categories and are common between the two sets. Next, we can recognize that each site (IMDB vs TMDB) have separate rating systems, with incongruencies in the number of votes and the average rating, though in this case the ratings are very close. It may work to choose the rating with the most votes at some point.

Well, now that we can easily access all this information, let's write a function to apply to our tmdb_movies dataframe and add columns with all the desirable information.

In [27]: `orig_dfs['tmdb_movies'].head()`
executed in 14ms, finished 13:40:07 2021-05-01

Out[27]:

	genre_ids	id	original_language	original_title	popularity	release_date	title	vote_aver
0	[12, 14, 10751]	12444	en	Harry Potter and the Deathly Hallows: Part 1	33.533	2010-11-19	Harry Potter and the Deathly Hallows: Part 1	7
1	[14, 12, 16, 10751]	10191	en	How to Train Your Dragon	28.734	2010-03-26	How to Train Your Dragon	7
2	[12, 28, 878]	10138	en	Iron Man 2	28.515	2010-05-07	Iron Man 2	6
3	[16, 35, 10751]	862	en	Toy Story	28.005	1995-11-22	Toy Story	7
4	[28, 878, 12]	27205	en	Inception	27.920	2010-07-16	Inception	8

```
In [28]: def grab_data_tmdb(row):
    movie_id = row['id']
    row_list = []
    try:
        movie_info = tmdb.Movies(movie_id).info()
    except:
        print('Missing or changed Movie ID')
        row_list = [np.nan] * 6
    else:
        row_list.extend( [ movie_info['runtime'],
                          movie_info['budget'],
                          movie_info['revenue'],
                          movie_info['imdb_id'] ] )
        if movie_info['title'] != row['original_title'] and movie_info['title']:
            print('Non-matching Movie Titles')
            row_list.append(1)
        else:
            row_list.append(0)
        production_raw = movie_info['production_companies']
        company_names = []
        for company in production_raw:
            company_names.append(company['name'])
        row_list.append(company_names)
    finally:
        return row.append(pd.Series(row_list, index = ['runtime', 'budget_est', 're']))
```

executed in 14ms, finished 13:40:07 2021-05-01

```
In [29]: # expanded_tmdb = orig_dfs['tmdb_movies'].apply(grab_data_tmdb, axis = 1)
```

executed in 15ms, finished 13:40:08 2021-05-01

```
In [30]: # expanded_tmdb.head()
```

executed in 14ms, finished 13:40:08 2021-05-01

Nice! Seems to have worked. Above, there was some output indicating missing or potentially changed movie IDs, so we will have to look into that. We will also have to see how many films didn't have data for the columns we hoped to retrieve. However, since it took an hour to collect the extra data via API requests, lets save the expanded dataframe as a new csv for future retrieval.

```
In [31]: # expanded_tmdb.to_csv('./dataFiles/expanded_tmdb.csv', index=False)
```

executed in 13ms, finished 13:40:08 2021-05-01

```
In [32]: expanded_tmdb = pd.read_csv('./dataFiles/expanded_tmdb.csv')
```

executed in 61ms, finished 13:40:09 2021-05-01

In [33]: `expanded_tmdb.head()`

executed in 14ms, finished 13:40:09 2021-05-01

Out[33]:

	genre_ids	id	original_language	original_title	popularity	release_date	title	vote_aver
0	[12, 14, 10751]	12444	en	Harry Potter and the Deathly Hallows: Part 1	33.533	2010-11-19	Harry Potter and the Deathly Hallows: Part 1	7
1	[14, 12, 16, 10751]	10191	en	How to Train Your Dragon	28.734	2010-03-26	How to Train Your Dragon	7
2	[12, 28, 878]	10138	en	Iron Man 2	28.515	2010-05-07	Iron Man 2	6
3	[16, 35, 10751]	862	en	Toy Story	28.005	1995-11-22	Toy Story	7
4	[28, 878, 12]	27205	en	Inception	27.920	2010-07-16	Inception	8

Now, let's see if we actually got any extra data. Since our budgetary information is most important, let's see how much of the original dataframe we were able to find this extra information for:

In [34]: `cleaned_tmdb = expanded_tmdb.loc[(expanded_tmdb['budget_est'] != 0) & (expanded_tmdb['revenue'] != 0) & (expanded_tmdb['req_title'] != 1)]`

executed in 12ms, finished 13:40:09 2021-05-01

In [35]: `cleaned_tmdb = cleaned_tmdb.dropna(axis = 0)
cleaned_tmdb = cleaned_tmdb.reset_index()
cleaned_tmdb.shape`

executed in 10ms, finished 13:40:10 2021-05-01

Out[35]: (2348, 16)

Okay.. not great. So, out of roughly 27000 movies, we were able to collect budget and revenue data for just under 10% of the records. Note: we also required that the title from the request response was the same as either the original_title or title column value from the original dataframe. This is to account for movie IDs that have been changed in the database since the original data file was retrieved, which is stated as possible in their API documentation. If we wanted to do extra work to check responses where the title did not match either, i.e. req_title = 1 in the expanded dataframe, we could potentially have an additional 450 records with financial information. At this time, it does not seem worth expending the effort, as anything up to date probably contains the

most popular and highest grossing films as is. So, in a roundabout way, all we have done is to reduce the tmdb_movies dataframe to a smaller size but with extra useful information, including (hopefully, if TMDB has properly sourced and updated this information) the ability to easily make connections to the imdb tables using toconst. Let's clean this table a little bit further below.

```
In [36]: cleaned_tmdb = cleaned_tmdb.drop(columns = ['index', 'req_title'])
```

executed in 14ms, finished 13:40:10 2021-05-01

```
In [37]: cleaned_tmdb = cleaned_tmdb.rename(columns = {'id': 'tmdb_id'})
```

executed in 14ms, finished 13:40:10 2021-05-01

Since our dream of a big dataset all from TMDB more or less failed without going further, we will get back to combining the tables we have with that information. As a reminder, lets .head() them here for quick reference.

```
In [38]: cleaned_tmdb.head()
```

executed in 16ms, finished 13:40:11 2021-05-01

Out[38]:

	genre_ids	tmdb_id	original_language	original_title	popularity	release_date	title	vote_ave
0	[12, 14, 10751]	12444	en	Harry Potter and the Deathly Hallows: Part 1	33.533	2010-11-19	Harry Potter and the Deathly Hallows: Part 1	
1	[14, 12, 16, 10751]	10191	en	How to Train Your Dragon	28.734	2010-03-26	How to Train Your Dragon	
2	[12, 28, 878]	10138	en	Iron Man 2	28.515	2010-05-07	Iron Man 2	
3	[16, 35, 10751]	862	en	Toy Story	28.005	1995-11-22	Toy Story	
4	[28, 878, 12]	27205	en	Inception	27.920	2010-07-16	Inception	

In [39]: `orig_dfs['bom_movie_gross'].head()`

executed in 15ms, finished 13:40:11 2021-05-01

Out[39]:

	title	studio	domestic_gross	foreign_gross	year
0	Toy Story 3	BV	415000000.0	652000000	2010
1	Alice in Wonderland (2010)	BV	334200000.0	691300000	2010
2	Harry Potter and the Deathly Hallows Part 1	WB	296000000.0	664300000	2010
3	Inception	WB	292600000.0	535700000	2010
4	Shrek Forever After	P/DW	238700000.0	513900000	2010

In [40]: `orig_dfs['tn_movie_budgets'].head()`

executed in 14ms, finished 13:40:11 2021-05-01

Out[40]:

	release_date	movie	production_budget	domestic_gross	worldwide_gross
0	Dec 18, 2009	Avatar	\$425,000,000	\$760,507,625	\$2,776,345,279
1	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	\$410,600,000	\$241,063,875	\$1,045,663,875
2	Jun 7, 2019	Dark Phoenix	\$350,000,000	\$42,762,350	\$149,762,350
3	May 1, 2015	Avengers: Age of Ultron	\$330,600,000	\$459,005,868	\$1,403,013,963
4	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	\$317,000,000	\$620,181,382	\$1,316,721,747

Lets double check the .info and some .value_counts() for each, to see if they're all ready for merging.

In [41]: `cleaned_tmdb.info()`

executed in 14ms, finished 13:40:12 2021-05-01

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2348 entries, 0 to 2347
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   genre_ids        2348 non-null    object  
 1   tmdb_id          2348 non-null    int64  
 2   original_language 2348 non-null    object  
 3   original_title    2348 non-null    object  
 4   popularity        2348 non-null    float64 
 5   release_date      2348 non-null    object  
 6   title             2348 non-null    object  
 7   vote_average      2348 non-null    float64 
 8   vote_count        2348 non-null    int64  
 9   runtime            2348 non-null    float64 
 10  budget_est        2348 non-null    float64 
 11  revenue            2348 non-null    float64 
 12  imdb_id           2348 non-null    object  
 13  production_companies 2348 non-null    object  
dtypes: float64(5), int64(2), object(7)
memory usage: 256.9+ KB
```

In [42]: `orig_dfs['bom_movie_gross'].info()`

executed in 14ms, finished 13:40:12 2021-05-01

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3387 entries, 0 to 3386
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   title             3387 non-null    object  
 1   studio            3382 non-null    object  
 2   domestic_gross    3359 non-null    float64 
 3   foreign_gross     2037 non-null    object  
 4   year              3387 non-null    int64  
dtypes: float64(1), int64(1), object(3)
memory usage: 132.4+ KB
```

```
In [43]: orig_dfs['bom_movie_gross'].foreign_gross.value_counts()
```

executed in 14ms, finished 13:40:12 2021-05-01

```
Out[43]: 1200000    23
1100000    14
4200000    12
1900000    12
2500000    11
...
46200000    1
128000000   1
93700000   1
29000000   1
63300000   1
Name: foreign_gross, Length: 1204, dtype: int64
```

```
In [44]: orig_dfs['tn_movie_budgets'].info()
```

executed in 14ms, finished 13:40:13 2021-05-01

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5782 entries, 0 to 5781
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   release_date     5782 non-null   object 
 1   movie            5782 non-null   object 
 2   production_budget 5782 non-null   object 
 3   domestic_gross   5782 non-null   object 
 4   worldwide_gross  5782 non-null   object 
dtypes: object(5)
memory usage: 226.0+ KB
```

```
In [45]: orig_dfs['tn_movie_budgets'].production_budget.value_counts()
```

executed in 14ms, finished 13:40:13 2021-05-01

```
Out[45]: $20,000,000    231
$10,000,000    212
$30,000,000    177
$15,000,000    173
$25,000,000    171
...
$6,600,000      1
$640,000        1
$270,000        1
$217,000,000    1
$730,000        1
Name: production_budget, Length: 509, dtype: int64
```

In [46]: `orig_dfs['tn_movie_budgets'].worldwide_gross.value_counts()`

executed in 15ms, finished 13:40:13 2021-05-01

Out[46]:

```
$0           367
$8,000,000      9
$7,000,000      6
$2,000,000      6
$11,000,000     4
...
$67,255,916      1
$10,447,579      1
$205,298,907     1
$13,205,411      1
$211,562,435     1
Name: worldwide_gross, Length: 5356, dtype: int64
```

Okay, so it looks like in the BOM data there are a lot of NULLs in foreign_gross, whereas for the TN data there are a lot of zero values, which could be either unreleased movies or movies for which this data hasn't been collected / reported. Additionally, all of the TN columns are object type, so we need to change the type of those columns to int64. Since we need at least this information for both, lets convert data types then get rid of any rows with NaNs or zeros in each dataframe before merging.

In [47]: `cleaned_bom = orig_dfs['bom_movie_gross'].dropna(axis = 0)`

`cleaned_bom = cleaned_bom.reset_index()`

`cleaned_bom = cleaned_bom.drop(columns = 'index')`

`cleaned_bom.info()`

executed in 15ms, finished 13:40:13 2021-05-01

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2007 entries, 0 to 2006
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   title            2007 non-null    object  
 1   studio           2007 non-null    object  
 2   domestic_gross   2007 non-null    float64
 3   foreign_gross    2007 non-null    object  
 4   year             2007 non-null    int64  
dtypes: float64(1), int64(1), object(3)
memory usage: 78.5+ KB
```

In [48]: `cleaned_bom.foreign_gross.value_counts()`

executed in 12ms, finished 13:40:14 2021-05-01

Out[48]:

1200000	23
1100000	13
4200000	12
2500000	11
1900000	11
	..
76000000	1
161000	1
45400000	1
133100000	1
63300000	1

Name: foreign_gross, Length: 1193, dtype: int64

In [49]: `cleaned_tn = orig_dfs['tn_movie_budgets'].copy()`

```
cleaned_tn['production_budget'] = cleaned_tn['production_budget'].apply(lambda x: int(x))
cleaned_tn['domestic_gross'] = cleaned_tn['domestic_gross'].apply(lambda x: int(x))
cleaned_tn['worldwide_gross'] = cleaned_tn['worldwide_gross'].apply(lambda x: int(x))
cleaned_tn.head()
```

executed in 31ms, finished 13:40:14 2021-05-01

Out[49]:

	release_date	movie	production_budget	domestic_gross	worldwide_gross
0	Dec 18, 2009	Avatar	425000000	760507625	2776345279
1	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	410600000	241063875	1045663875
2	Jun 7, 2019	Dark Phoenix	350000000	42762350	149762350
3	May 1, 2015	Avengers: Age of Ultron	330600000	459005868	1403013963
4	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	317000000	620181382	1316721747

In [50]: `cleaned_tn = cleaned_tn.loc[(cleaned_tn['production_budget'] != 0) &`

```
(cleaned_tn['domestic_gross'] != 0) &
(cleaned_tn['worldwide_gross'] != 0) ]
```

executed in 14ms, finished 13:40:15 2021-05-01

In [51]: `cleaned_tn.production_budget.value_counts()`

executed in 10ms, finished 13:40:15 2021-05-01

Out[51]: 20000000 221

10000000 197

30000000 171

40000000 163

15000000 162

...

420000 1

7303082 1

306000000 1

4638783 1

28500000 1

Name: production_budget, Length: 472, dtype: int64

In [52]: `cleaned_tn.worldwide_gross.value_counts()`

executed in 14ms, finished 13:40:16 2021-05-01

Out[52]: 8000000 9

2000000 6

7000000 6

11000000 4

9000000 4

..

478595 1

58545540 1

92618117 1

3902679 1

104267443 1

Name: worldwide_gross, Length: 5177, dtype: int64

In [53]: `cleaned_tn = cleaned_tn.reset_index()`

`cleaned_tn = cleaned_tn.drop(columns = 'index')`

`cleaned_tn.head()`

executed in 15ms, finished 13:40:16 2021-05-01

Out[53]:

	release_date	movie	production_budget	domestic_gross	worldwide_gross
0	Dec 18, 2009	Avatar	425000000	760507625	2776345279
1	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	410600000	241063875	1045663875
2	Jun 7, 2019	Dark Phoenix	350000000	42762350	149762350
3	May 1, 2015	Avengers: Age of Ultron	330600000	459005868	1403013963
4	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	317000000	620181382	1316721747

```
In [54]: print(cleaned_tmdb.tmdb_id.nunique())
print(cleaned_tn.movie.nunique())
print(cleaned_tmdb.shape)
print(cleaned_tn.shape)
```

executed in 26ms, finished 13:40:16 2021-05-01

```
2143
5159
(2348, 14)
(5234, 5)
```

Hmm... it looks like there may be some duplicate movie records to take care of. A couple hundred in each, let's remove them.

```
In [55]: cleaned_tmdb = cleaned_tmdb.drop_duplicates(subset = ['tmdb_id'])
cleaned_tn = cleaned_tn.drop_duplicates(subset = ['movie', 'release_date'])
```

executed in 14ms, finished 13:40:17 2021-05-01

```
In [56]: print(cleaned_tmdb.tmdb_id.nunique())
print(cleaned_tn.movie.nunique())
print(cleaned_tmdb.shape)
print(cleaned_tn.shape)
```

executed in 13ms, finished 13:40:17 2021-05-01

```
2143
5159
(2143, 14)
(5234, 5)
```

I also am noticing that there may be some duplicate movie titles with different release dates that we want to keep as separate records. The easiest way to handle this will be to separate out the release date into three columns each, release day, month, and year, then finally merge on both the movie title and the release year.

```
In [57]: def create_date_cols(row):
    timestamp_obj = pd.to_datetime(row['release_date'])
    day = timestamp_obj.day
    month = timestamp_obj.month
    year = timestamp_obj.year
    day_of_week = timestamp_obj.dayofweek
    return row.append(pd.Series([day_of_week, day, month, year],
                               index=['release_week_day', 'release_day', 'releas
```

executed in 13ms, finished 13:40:18 2021-05-01

```
In [58]: cleaned_tmdb = cleaned_tmdb.apply(create_date_cols, axis=1)
cleaned_tn = cleaned_tn.apply(create_date_cols, axis=1)
```

executed in 4.27s, finished 13:40:22 2021-05-01

In [59]: `cleaned_tn.head()`

executed in 14ms, finished 13:40:23 2021-05-01

Out[59]:

	release_date	movie	production_budget	domestic_gross	worldwide_gross	release_week_day
0	Dec 18, 2009	Avatar	425000000	760507625	2776345279	
1	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	410600000	241063875	1045663875	
2	Jun 7, 2019	Dark Phoenix	350000000	42762350	149762350	
3	May 1, 2015	Avengers: Age of Ultron	330600000	459005868	1403013963	
4	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	317000000	620181382	1316721747	

In [60]: `cleaned_tmdb = cleaned_tmdb.drop(columns = 'release_date')`
`cleaned_tn = cleaned_tn.drop(columns = 'release_date')`

executed in 15ms, finished 13:40:23 2021-05-01

I'm thinking now that I would like to see what other information I can add to the TN dataset before merging. For example, we can use title_basics to add a tconst (better matching), the runtime, and genres. Then, since that table will have tconst, we can add information from the title_ratings dataframe. So, cleaned_tn <- title_basics <- title_ratings. Let's check them out pre-merging:

In [61]: `orig_dfs['title_basics'].head()`

executed in 14ms, finished 13:40:23 2021-05-01

Out[61]:

	tconst	primary_title	original_title	start_year	runtime_minutes	genres
0	tt0063540	Sunghursh	Sunghursh	2013	175.0	Action,Crime,Drama
1	tt0066787	One Day Before the Rainy Season	Ashad Ka Ek Din	2019	114.0	Biography,Drama
2	tt0069049	The Other Side of the Wind	The Other Side of the Wind	2018	122.0	Drama
3	tt0069204	Sabse Bada Sukh	Sabse Bada Sukh	2018	NaN	Comedy,Drama
4	tt0100275	The Wandering Soap Opera	La Telenovela Errante	2017	80.0	Comedy,Drama,Fantasy

In [62]: `orig_dfs['title_basics'].info()`

executed in 25ms, finished 13:40:24 2021-05-01

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 146144 entries, 0 to 146143
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   tconst          146144 non-null   object  
 1   primary_title    146144 non-null   object  
 2   original_title   146123 non-null   object  
 3   start_year       146144 non-null   int64   
 4   runtime_minutes  114405 non-null   float64 
 5   genres           140736 non-null   object  
dtypes: float64(1), int64(1), object(4)
memory usage: 6.7+ MB
```

In [63]: `orig_dfs['title_basics'].tconst.nunique()`

executed in 30ms, finished 13:40:24 2021-05-01

Out[63]: 146144

Okay, so there are some missing values here and there, but the ones we will want to join on, tconst, titles and start_year are basically all there, and there are no duplicate rows based on tconst.

In [64]: `orig_dfs['title_ratings'].head()`

executed in 15ms, finished 13:40:25 2021-05-01

Out[64]:

	tconst	averagerating	numvotes
0	tt10356526	8.3	31
1	tt10384606	8.9	559
2	tt1042974	6.4	20
3	tt1043726	4.2	50352
4	tt1060240	6.5	21

In [65]: `orig_dfs['title_ratings'].info()`

executed in 15ms, finished 13:40:25 2021-05-01

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 73856 entries, 0 to 73855
Data columns (total 3 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   tconst          73856 non-null   object  
 1   averagerating   73856 non-null   float64 
 2   numvotes        73856 non-null   int64  
dtypes: float64(1), int64(1), object(1)
memory usage: 1.7+ MB
```

```
In [66]: orig_dfs['title_ratings'].tconst.nunique()
```

executed in 15ms, finished 13:40:26 2021-05-01

```
Out[66]: 73856
```

The same seems to be true for title_ratings as well, though we have ratings information for many less movies than we have title basics for. Anyway, lets quickly merge the two title dataframes:

```
In [67]: titles_plus_ratings = orig_dfs['title_basics'].merge(orig_dfs['title_ratings'],
                                                       how='left',
                                                       on='tconst')
titles_plus_ratings.head(20)
```

executed in 92ms, finished 13:40:26 2021-05-01

Out[67]:

	tconst	primary_title	original_title	start_year	runtime_minutes	genres
0	tt0063540	Sunghursh	Sunghursh	2013	175.0	Action,Crime,Drama
1	tt0066787	One Day Before the Rainy Season	Ashad Ka Ek Din	2019	114.0	Biography,Drama
2	tt0069049	The Other Side of the Wind	The Other Side of the Wind	2018	122.0	Drama
3	tt0069204	Sabse Bada Sukh	Sabse Bada Sukh	2018	NaN	Comedy,Drama
4	tt0100275	The Wandering Soap Opera	La Telenovela Errante	2017	80.0	Comedy,Drama,Fantasy
5	tt0111414	A Thin Life	A Thin Life	2018	75.0	Comedy
6	tt0112502	Bigfoot	Bigfoot	2017	NaN	Horror,Thriller
7	tt0137204	Joe Finds Grace	Joe Finds Grace	2017	83.0	Adventure,Animation,Comedy
8	tt0139613	O Silêncio	O Silêncio	2012	NaN	Documentary,History
9	tt0144449	Nema aviona za Zagreb	Nema aviona za Zagreb	2012	82.0	Biography
10	tt0146592	Pál Adrienn	Pál Adrienn	2010	136.0	Drama
11	tt0154039	So Much for Justice!	Oda az igazság	2010	100.0	History
12	tt0159369	Cooper and Hemingway: The True Gen	Cooper and Hemingway: The True Gen	2013	180.0	Documentary
13	tt0162942	Children of the Green Dragon	A zöld sárkány gyermekai	2010	89.0	Drama
14	tt0170651	T.G.M. - osvoboditel	T.G.M. - osvoboditel	2018	60.0	Documentary
15	tt0176694	The Tragedy of Man	Az ember tragédiája	2011	160.0	Animation,Drama,History
16	tt0187902	How Huang Fei-hong Rescued the Orphan from the...	How Huang Fei-hong Rescued the Orphan from the...	2011	NaN	NaN
17	tt0192528	Heaven & Hell	Reverse Heaven	2018	104.0	Drama
18	tt0230212	The Final Journey	The Final Journey	2010	120.0	Drama

	tconst	primary_title	original_title	start_year	runtime_minutes	genres
19	tt0247643	Los pájaros se van con la muerte	Los pájaros se van con la muerte	2011	110.0	Drama,Mystery

In [68]: `titles_plus_ratings.shape[0] - titles_plus_ratings.dropna().shape[0]`

executed in 45ms, finished 13:40:26 2021-05-01

Out[68]: 80424

In [69]: `titles_plus_ratings.tconst.nunique()`

executed in 30ms, finished 13:40:27 2021-05-01

Out[69]: 146144

So, there are a majority of rows without complete information, but when we merge with cleaned_tn it may not matter.

In [70]: `cleaned_tn.shape`

executed in 14ms, finished 13:40:27 2021-05-01

Out[70]: (5234, 8)

Since there are multiple title type columns to match on, we can avoid duplicates by creating two dataframes, one where the original_title and primary_title are the same and one where they are different, and then join cleaned_tn to those before concatenating all the results.

In [71]: `diff_titles = titles_plus_ratings.loc[titles_plus_ratings.primary_title != titles_plus_ratings.original_title]`
`same_titles = titles_plus_ratings.loc[titles_plus_ratings.primary_title == titles_plus_ratings.original_title]`

executed in 43ms, finished 13:40:28 2021-05-01

In [72]: `diff_tn1 = cleaned_tn.merge(diff_titles,`
 `how='inner',`
 `left_on=['movie', 'release_year'],`
 `right_on=['primary_title', 'start_year'])`
`diff_tn2 = cleaned_tn.merge(diff_titles,`
 `how='inner',`
 `left_on=['movie', 'release_year'],`
 `right_on=['original_title', 'start_year'])`
`same_tn1 = cleaned_tn.merge(same_titles,`
 `how='inner',`
 `left_on=['movie', 'release_year'],`
 `right_on=['original_title', 'start_year'])`

executed in 59ms, finished 13:40:28 2021-05-01

```
In [73]: diff_tn1.shape
```

executed in 8ms, finished 13:40:42 2021-05-01

```
Out[73]: (32, 16)
```

```
In [74]: diff_tn2.shape
```

executed in 14ms, finished 13:40:43 2021-05-01

```
Out[74]: (18, 16)
```

```
In [75]: same_tn1.shape
```

executed in 14ms, finished 13:40:44 2021-05-01

```
Out[75]: (1338, 16)
```

```
In [76]: buffed_tn = pd.concat([diff_tn1,diff_tn2,same_tn1])
```

executed in 14ms, finished 13:40:44 2021-05-01

```
In [77]: buffed_tn.head()
```

executed in 14ms, finished 13:40:45 2021-05-01

```
Out[77]:
```

	movie	production_budget	domestic_gross	worldwide_gross	release_week_day	release_day	r
0	Rogue One: A Star Wars Story	2000000000	532177324	1049102856	4	16	
1	Iron Man 3	2000000000	408992272	1215392272	4	3	
2	Furious 7	1900000000	353007020	1518722794	4	3	
3	X-Men: First Class	1600000000	146408305	355408305	4	3	
4	Home	1300000000	177397510	385997896	4	27	

◀ ▶

```
In [78]: buffed_tn.shape
```

executed in 14ms, finished 13:40:45 2021-05-01

```
Out[78]: (1388, 16)
```

In [79]: `buffed_tn.info()`

executed in 14ms, finished 13:40:46 2021-05-01

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1388 entries, 0 to 1337
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   movie            1388 non-null    object  
 1   production_budget 1388 non-null    int64   
 2   domestic_gross    1388 non-null    int64   
 3   worldwide_gross   1388 non-null    int64   
 4   release_week_day  1388 non-null    int64   
 5   release_day       1388 non-null    int64   
 6   release_month     1388 non-null    int64   
 7   release_year      1388 non-null    int64   
 8   tconst            1388 non-null    object  
 9   primary_title     1388 non-null    object  
 10  original_title    1388 non-null    object  
 11  start_year        1388 non-null    int64   
 12  runtime_minutes   1375 non-null    float64 
 13  genres            1385 non-null    object  
 14  averagerating     1360 non-null    float64 
 15  numvotes          1360 non-null    float64 
dtypes: float64(3), int64(8), object(5)
memory usage: 184.3+ KB
```

Okay, we had a few rows sneak in from `titles_plus_ratings` where we had nulls for some of the values, so let's ditch those few rows.

```
In [80]: buffed_tn = buffed_tn.dropna(axis = 0)
buffed_tn = buffed_tn.reset_index()
buffed_tn = buffed_tn.drop(columns = 'index')
buffed_tn.info()
```

executed in 14ms, finished 13:40:46 2021-05-01

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1355 entries, 0 to 1354
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   movie            1355 non-null    object  
 1   production_budget 1355 non-null    int64   
 2   domestic_gross    1355 non-null    int64   
 3   worldwide_gross   1355 non-null    int64   
 4   release_week_day  1355 non-null    int64   
 5   release_day       1355 non-null    int64   
 6   release_month     1355 non-null    int64   
 7   release_year      1355 non-null    int64   
 8   tconst            1355 non-null    object  
 9   primary_title     1355 non-null    object  
 10  original_title    1355 non-null    object  
 11  start_year       1355 non-null    int64   
 12  runtime_minutes   1355 non-null    float64 
 13  genres            1355 non-null    object  
 14  averagerating     1355 non-null    float64 
 15  numvotes          1355 non-null    float64 
dtypes: float64(3), int64(8), object(5)
memory usage: 169.5+ KB
```

```
In [81]: buffed_tn.head()
```

executed in 14ms, finished 13:40:47 2021-05-01

Out[81]:

		movie	production_budget	domestic_gross	worldwide_gross	release_week_day	release_day	...
0	Rogue One: A Star Wars Story	200000000	532177324	1049102856	4	16		
1	Iron Man 3	200000000	408992272	1215392272	4	3		
2	Furious 7	190000000	353007020	1518722794	4	3		
3	X-Men: First Class	160000000	146408305	355408305	4	3		
4	The Promise	90000000	8224288	10551417	4	21		

Okay, so now we finally have two dataframes with as much categorical and numerical information for their respective movies as is currently possible, buffed_tn and cleaned_tmdb. We should be able to easily concatenate these and then remove duplicate rows based on tconst. Let's make them

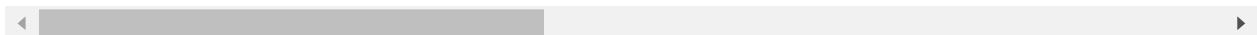
both contain all the same columns in the same order to prepare for concatenating.

In [82]: `cleaned_tmdb.head()`

executed in 14ms, finished 13:40:48 2021-05-01

Out[82]:

	genre_ids	tmdb_id	original_language	original_title	popularity	title	vote_average	vote_count
0	[12, 14, 10751]	12444	en	Harry Potter and the Deathly Hallows: Part 1	33.533	Harry Potter and the Deathly Hallows: Part 1	7.7	103
1	[14, 12, 16, 10751]	10191	en	How to Train Your Dragon	28.734	How to Train Your Dragon	7.7	76
2	[12, 28, 878]	10138	en	Iron Man 2	28.515	Iron Man 2	6.8	123
3	[16, 35, 10751]	862	en	Toy Story	28.005	Toy Story	7.9	103
4	[28, 878, 12]	27205	en	Inception	27.920	Inception	8.3	223



In [83]: `cleaned_tmdb.columns`

executed in 15ms, finished 13:40:48 2021-05-01

Out[83]: `Index(['genre_ids', 'tmdb_id', 'original_language', 'original_title', 'popularity', 'title', 'vote_average', 'vote_count', 'runtime', 'budget_est', 'revenue', 'imdb_id', 'production_companies', 'release_week_day', 'release_day', 'release_month', 'release_year'], dtype='object')`

In [84]: `cleaned_tmdb = cleaned_tmdb.drop(columns = ['tmdb_id', 'original_language', 'original_title'])`

executed in 14ms, finished 13:40:49 2021-05-01

In [85]: `cleaned_tmdb.columns`

executed in 15ms, finished 13:40:49 2021-05-01

Out[85]: `Index(['genre_ids', 'title', 'vote_average', 'vote_count', 'runtime', 'budget_est', 'revenue', 'imdb_id', 'release_week_day', 'release_day', 'release_month', 'release_year'], dtype='object')`

```
In [86]: buffed_tn.columns
```

executed in 15ms, finished 13:40:50 2021-05-01

```
Out[86]: Index(['movie', 'production_budget', 'domestic_gross', 'worldwide_gross',
       'release_week_day', 'release_day', 'release_month', 'release_year',
       'tconst', 'primary_title', 'original_title', 'start_year',
       'runtime_minutes', 'genres', 'averagerating', 'numvotes'],
      dtype='object')
```

```
In [87]: buffed_tn = buffed_tn.drop(columns = ['primary_title','original_title','domestic_
```

executed in 15ms, finished 13:40:50 2021-05-01

```
In [88]: buffed_tn.columns
```

executed in 14ms, finished 13:40:51 2021-05-01

```
Out[88]: Index(['movie', 'production_budget', 'worldwide_gross', 'release_week_day',
       'release_day', 'release_month', 'release_year', 'tconst',
       'runtime_minutes', 'genres', 'averagerating', 'numvotes'],
      dtype='object')
```

```
In [89]: cleaned_tmdb = cleaned_tmdb.rename(columns={'genre_ids': 'Genres', 'title': 'Tit
```

'runtime': 'runtime_minutes', 'revenue': 'worldwide_gross',

executed in 15ms, finished 13:40:51 2021-05-01

```
In [90]: cleaned_tmdb.columns
```

executed in 15ms, finished 13:40:52 2021-05-01

```
Out[90]: Index(['Genres', 'Title', 'vote_average', 'vote_count', 'runtime_minutes',
       'budget_est', 'worldwide_gross', 'tconst', 'release_week_day',
       'release_day', 'release_month', 'release_year'],
      dtype='object')
```

```
In [91]: buffed_tn = buffed_tn.rename(
```

columns={'movie': 'Title', 'production_budget': 'budget_est', 'genres': 'Genre',
'averagerating': 'vote_average', 'numvotes': 'vote_count'})

executed in 14ms, finished 13:40:53 2021-05-01

```
In [92]: buffed_tn.columns
```

executed in 14ms, finished 13:40:53 2021-05-01

```
Out[92]: Index(['Title', 'budget_est', 'worldwide_gross', 'release_week_day',
       'release_day', 'release_month', 'release_year', 'tconst',
       'runtime_minutes', 'Genres', 'vote_average', 'vote_count'],
      dtype='object')
```

```
In [93]: len(buffed_tn.columns)
```

executed in 14ms, finished 13:40:54 2021-05-01

```
Out[93]: 12
```

In [94]: `len(cleaned_tmdb.columns)`

executed in 15ms, finished 13:40:54 2021-05-01

Out[94]: 12

In [95]: `cleaned_tmdb = cleaned_tmdb[['tconst', 'Title', 'Genres', 'runtime_minutes', 'vote_average', 'vote_count', 'budget_est', 'worldwide_gross', 'release_week_day', 'release_day', 'release_month']]`

executed in 14ms, finished 13:40:55 2021-05-01

In [96]: `cleaned_tmdb.head()`

executed in 28ms, finished 13:40:55 2021-05-01

Out[96]:

	tconst	Title	Genres	runtime_minutes	vote_average	vote_count	budget_est	worldwide_gross
0	tt0926084	Harry Potter and the Deathly Hallows: Part 1	[12, 14, 10751]	146.0	7.7	10788	250000000.0	9543
1	tt0892769	How to Train Your Dragon	[14, 12, 16, 10751]	98.0	7.7	7610	165000000.0	4948
2	tt1228705	Iron Man 2	[12, 28, 878]	124.0	6.8	12368	200000000.0	6239
3	tt0114709	Toy Story	[16, 35, 10751]	81.0	7.9	10174	30000000.0	3735
4	tt1375666	Inception	[28, 878, 12]	148.0	8.3	22186	160000000.0	8255

In [97]: `buffed_tn = buffed_tn[['tconst', 'Title', 'Genres', 'runtime_minutes', 'vote_average', 'vote_count', 'budget_est', 'worldwide_gross', 'release_week_day', 'release_day', 'release_month']]`

executed in 15ms, finished 13:40:56 2021-05-01

In [98]: `buffed_tn.head()`

executed in 27ms, finished 13:40:56 2021-05-01

Out[98]:

	tconst	Title	Genres	runtime_minutes	vote_average	vote_count	budget_est
0	tt3748528	Rogue One: A Star Wars Story	Action,Adventure,Sci-Fi	133.0	7.8	478592.0	2000000000
1	tt1300854	Iron Man 3	Action,Adventure,Sci-Fi	130.0	7.2	692794.0	2000000000
2	tt2820852	Furious 7	Action,Crime,Thriller	137.0	7.2	335074.0	1900000000
3	tt1270798	X-Men: First Class	Action,Adventure,Sci-Fi	131.0	7.7	608930.0	1600000000
4	tt7232438	The Promise	Drama,Horror,Thriller	114.0	6.1	629.0	900000000

There is one last thing to do before we can work with these. We need to fix that in `buffed_tn`. `Genres` is a comma separated string and in `cleaned_tmdb` it is a list of integer IDs. For `buffed_tn`, we need to simply turn the string into a list using a comma as a delimiter. For `cleaned_tmdb` we have a dictionary that we can convert the integer ID's into string values for, and then leave as a list. Thus a list of genres as strings in both dataframes.

In [99]: `buffed_tn['Genres'] = buffed_tn['Genres'].apply(lambda x: x.split(','))`

executed in 13ms, finished 13:40:57 2021-05-01

In [100]: `buffed_tn.head()`

executed in 15ms, finished 13:40:58 2021-05-01

Out[100]:

	tconst	Title	Genres	runtime_minutes	vote_average	vote_count	budget_est	worldwid
0	tt3748528	Rogue One: A Star Wars Story	[Action, Adventure, Sci-Fi]	133.0	7.8	478592.0	200000000	104
1	tt1300854	Iron Man 3	[Action, Adventure, Sci-Fi]	130.0	7.2	692794.0	200000000	121
2	tt2820852	Furious 7	[Action, Crime, Thriller]	137.0	7.2	335074.0	190000000	151
3	tt1270798	X-Men: First Class	[Action, Adventure, Sci-Fi]	131.0	7.7	608930.0	160000000	35
4	tt7232438	The Promise	[Drama, Horror, Thriller]	114.0	6.1	629.0	90000000	1



In [101]: `buffed_tn.Genres[0]`

executed in 14ms, finished 13:40:58 2021-05-01

Out[101]: `['Action', 'Adventure', 'Sci-Fi']`

For cleaned_tmdb, it's a bit more challenging because even though it looks like a list, it is really stored as a string, e.g. '[1,2,3]' instead of [1,2,3]. I found this module to handle that, ast, and it will convert the former to the latter.

In [102]: `import ast`

executed in 11ms, finished 13:40:59 2021-05-01

In [103]: `ast.literal_eval(cleaned_tmdb.Genres[0])`

executed in 14ms, finished 13:40:59 2021-05-01

Out[103]: `[12, 14, 10751]`

```
In [104]: genre_dict = {28: 'Action',
                     12: 'Adventure',
                     16: 'Animation',
                     35: 'Comedy',
                     80: 'Crime',
                     99: 'Documentary',
                     18: 'Drama',
                     10751: 'Family',
                     14: 'Fantasy',
                     36: 'History',
                     27: 'Horror',
                     10402: 'Music',
                     9648: 'Mystery',
                     10749: 'Romance',
                     878: 'Science Fiction',
                     10770: 'TV Movie',
                     53: 'Thriller',
                     10752: 'War',
                     37: 'Western'}
cleaned_tmdb['Genres'] = cleaned_tmdb['Genres'].apply(
    lambda x: [genre_dict[genre_id] for genre_id in ast.literal_eval(x)])
```

executed in 44ms, finished 13:41:00 2021-05-01

```
In [105]: cleaned_tmdb.Genres[0]
```

executed in 15ms, finished 13:41:01 2021-05-01

```
Out[105]: ['Adventure', 'Fantasy', 'Family']
```

In [106]: `cleaned_tmdb.head()`

executed in 15ms, finished 13:41:01 2021-05-01

Out[106]:

	tconst	Title	Genres	runtime_minutes	vote_average	vote_count	budget_est	worldwide_gross
0	tt0926084	Harry Potter and the Deathly Hallows: Part 1	[Adventure, Fantasy, Family]	146.0	7.7	10788	250000000.0	9
1	tt0892769	How to Train Your Dragon	[Fantasy, Adventure, Animation, Family]	98.0	7.7	7610	165000000.0	4
2	tt1228705	Iron Man 2	[Adventure, Action, Science Fiction]	124.0	6.8	12368	200000000.0	6
3	tt0114709	Toy Story	[Animation, Comedy, Family]	81.0	7.9	10174	30000000.0	3
4	tt1375666	Inception	[Action, Science Fiction, Adventure]	148.0	8.3	22186	160000000.0	8

Finally, let's concatenate these two dataframes. It is likely there are some duplicate movies, so we will have to weed them out. For the duplicates, however, instead of just taking the first and dropping the second, since they are from different sources and we are using a combined dataset, I would like to average things like the budget est and worldwide gross, and do a weighted reaveraging of the vote_average and counts. So, let's merge, find the subsection of duplicates and deal with them using a second merge, and then find any entries not in the intersection from each original dataframe and include those records as well.

In [107]: `movie_collection_w_dups = pd.concat([cleaned_tmdb, buffed_tn])`

executed in 12ms, finished 13:41:02 2021-05-01

In [108]: `movie_collection_w_dups.head()`

executed in 27ms, finished 13:41:02 2021-05-01

Out[108]:

	tconst	Title	Genres	runtime_minutes	vote_average	vote_count	budget_est	worldwide_gross
0	tt0926084	Harry Potter and the Deathly Hallows: Part 1	[Adventure, Fantasy, Family]	146.0	7.7	10788.0	250000000.0	9
1	tt0892769	How to Train Your Dragon	[Fantasy, Adventure, Animation, Family]	98.0	7.7	7610.0	165000000.0	4
2	tt1228705	Iron Man 2	[Adventure, Action, Science Fiction]	124.0	6.8	12368.0	200000000.0	6
3	tt0114709	Toy Story	[Animation, Comedy, Family]	81.0	7.9	10174.0	30000000.0	3
4	tt1375666	Inception	[Action, Science Fiction, Adventure]	148.0	8.3	22186.0	160000000.0	8

In [109]: `movie_collection_w_dups.info()`

executed in 15ms, finished 13:41:03 2021-05-01

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3498 entries, 0 to 1354
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   tconst          3498 non-null    object 
 1   Title            3498 non-null    object 
 2   Genres           3498 non-null    object 
 3   runtime_minutes  3498 non-null    float64
 4   vote_average     3498 non-null    float64
 5   vote_count       3498 non-null    float64
 6   budget_est       3498 non-null    float64
 7   worldwide_gross  3498 non-null    float64
 8   release_week_day 3498 non-null    int64  
 9   release_day      3498 non-null    int64  
 10  release_month    3498 non-null    int64  
 11  release_year     3498 non-null    int64  
dtypes: float64(5), int64(4), object(3)
memory usage: 355.3+ KB
```

```
In [110]: movie_collection_w_dups.tconst.nunique()
```

executed in 14ms, finished 13:41:03 2021-05-01

```
Out[110]: 2373
```

So, out of ~3500 records there are only 2373 unique movies.

```
In [111]: inner_sect = cleaned_tmdb.merge(buffed_tn, on = 'tconst')
```

executed in 15ms, finished 13:41:04 2021-05-01

```
In [112]: inner_sect.info()
```

executed in 15ms, finished 13:41:05 2021-05-01

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1125 entries, 0 to 1124
Data columns (total 23 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   tconst            1125 non-null   object  
 1   Title_x           1125 non-null   object  
 2   Genres_x          1125 non-null   object  
 3   runtime_minutes_x 1125 non-null   float64 
 4   vote_average_x    1125 non-null   float64 
 5   vote_count_x      1125 non-null   int64   
 6   budget_est_x      1125 non-null   float64 
 7   worldwide_gross_x 1125 non-null   float64 
 8   release_week_day_x 1125 non-null   int64   
 9   release_day_x     1125 non-null   int64   
 10  release_month_x   1125 non-null   int64   
 11  release_year_x    1125 non-null   int64   
 12  Title_y           1125 non-null   object  
 13  Genres_y          1125 non-null   object  
 14  runtime_minutes_y 1125 non-null   float64 
 15  vote_average_y    1125 non-null   float64 
 16  vote_count_y      1125 non-null   float64 
 17  budget_est_y      1125 non-null   int64   
 18  worldwide_gross_y 1125 non-null   int64   
 19  release_week_day_y 1125 non-null   int64   
 20  release_day_y     1125 non-null   int64   
 21  release_month_y   1125 non-null   int64   
 22  release_year_y    1125 non-null   int64   
dtypes: float64(7), int64(11), object(5)
memory usage: 210.9+ KB
```

In [113]: `inner_sect.head()`

executed in 30ms, finished 13:41:05 2021-05-01

Out[113]:

	tconst	Title_x	Genres_x	runtime_minutes_x	vote_average_x	vote_count_x	budget_est
0	tt0892769	How to Train Your Dragon	[Fantasy, Adventure, Animation, Family]	98.0	7.7	7610	165000000C
1	tt1228705	Iron Man 2	[Adventure, Action, Science Fiction]	124.0	6.8	12368	200000000C
2	tt1375666	Inception	[Action, Science Fiction, Adventure]	148.0	8.3	22186	160000000C
3	tt0814255	Percy Jackson & the Olympians: The Lightning T...	[Adventure, Fantasy, Family]	118.0	6.1	4229	95000000C
4	tt0435761	Toy Story 3	[Animation, Family, Comedy]	103.0	7.7	8340	200000000C

5 rows × 23 columns

In [114]: `common_tconst = pd.DataFrame([inner_sect.tconst,`

```
inner_sect.Title_x,
inner_sect.Genres_x,
(inner_sect.runtime_minutes_x + inner_sect.runtime_
((inner_sect.vote_average_x * inner_sect.vote_count_
(inner_sect.vote_count_x + inner_sect.vote_count_y),
(inner_sect.budget_est_x + inner_sect.budget_est_y),
(inner_sect.worldwide_gross_x + inner_sect.worldwide_
inner_sect.release_week_day_x, inner_sect.release_d
inner_sect.release_month_x, inner_sect.release_year
```

executed in 92ms, finished 13:41:06 2021-05-01

In [115]: `common_tconst.head()`

executed in 15ms, finished 13:41:07 2021-05-01

Out[115]:

	tconst	Title_x	Genres_x	Unnamed_0	Unnamed_1	Unnamed_2	Unnamed_3	Unnamed_4	rel
0	tt0892769	How to Train Your Dragon	[Fantasy, Adventure, Animation, Family]	98	8.09508	618909	1.65e+08	4.94875e+08	
1	tt1228705	Iron Man 2	[Adventure, Action, Science Fiction]	124	6.99631	670058	1.85e+08	6.22545e+08	
2	tt1375666	Inception	[Action, Science Fiction, Adventure]	148	8.79405	1.86325e+06	1.6e+08	8.30529e+08	
3	tt0814255	Percy Jackson & the Olympians: The Lightning T...	[Adventure, Fantasy, Family]	118	5.90499	169337	9.5e+07	2.24774e+08	
4	tt0435761	Toy Story 3	[Animation, Family, Comedy]	103	8.29275	690558	2e+08	1.06792e+09	

In [116]: `common_tconst.columns = cleaned_tmdb.columns`

executed in 14ms, finished 13:41:07 2021-05-01

In [117]: `common_tconst.head()`

executed in 15ms, finished 13:41:08 2021-05-01

Out[117]:

	tconst	Title	Genres	runtime_minutes	vote_average	vote_count	budget_est	wor...
0	tt0892769	How to Train Your Dragon	[Fantasy, Adventure, Animation, Family]	98	8.09508	618909	1.65e+08	
1	tt1228705	Iron Man 2	[Adventure, Action, Science Fiction]	124	6.99631	670058	1.85e+08	
2	tt1375666	Inception	[Action, Science Fiction, Adventure]	148	8.79405	1.86325e+06	1.6e+08	
3	tt0814255	Percy Jackson & the Olympians: The Lightning T...	[Adventure, Fantasy, Family]	118	5.90499	169337	9.5e+07	
4	tt0435761	Toy Story 3	[Animation, Family, Comedy]	103	8.29275	690558	2e+08	

In [118]: `outer_sect = movie_collection_w_dups.loc[~movie_collection_w_dups.tconst.isin(common_tconst)]`

executed in 15ms, finished 13:41:09 2021-05-01

In [119]: `outer_sect.head()`

executed in 13ms, finished 13:41:09 2021-05-01

Out[119]:

	tconst	Title	Genres	runtime_minutes	vote_average	vote_count	budget_est	wo
0	tt0926084	Harry Potter and the Deathly Hallows: Part 1	[Adventure, Fantasy, Family]	146.0	7.7	10788.0	250000000.0	
3	tt0114709	Toy Story	[Animation, Comedy, Family]	81.0	7.9	10174.0	30000000.0	
6	tt0499549	Avatar	[Action, Adventure, Fantasy, Science Fiction]	162.0	7.4	18676.0	237000000.0	
10	tt0120363	Toy Story 2	[Animation, Comedy, Family]	92.0	7.5	7553.0	90000000.0	
17	tt0372183	The Bourne Supremacy	[Action, Drama, Thriller]	108.0	7.3	4367.0	75000000.0	

In [120]: `outer_sect.info()`

executed in 11ms, finished 13:41:10 2021-05-01

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1248 entries, 0 to 1354
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   tconst            1248 non-null   object  
 1   Title             1248 non-null   object  
 2   Genres            1248 non-null   object  
 3   runtime_minutes   1248 non-null   float64 
 4   vote_average      1248 non-null   float64 
 5   vote_count        1248 non-null   float64 
 6   budget_est        1248 non-null   float64 
 7   worldwide_gross   1248 non-null   float64 
 8   release_week_day  1248 non-null   int64  
 9   release_day       1248 non-null   int64  
 10  release_month     1248 non-null   int64  
 11  release_year      1248 non-null   int64  
dtypes: float64(5), int64(4), object(3)
memory usage: 126.8+ KB
```

In [121]: `common_tconst.info()`

executed in 14ms, finished 13:41:11 2021-05-01

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1125 entries, 0 to 1124
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   tconst            1125 non-null   object  
 1   Title             1125 non-null   object  
 2   Genres            1125 non-null   object  
 3   runtime_minutes   1125 non-null   object  
 4   vote_average      1125 non-null   object  
 5   vote_count        1125 non-null   object  
 6   budget_est        1125 non-null   object  
 7   worldwide_gross   1125 non-null   object  
 8   release_week_day 1125 non-null   object  
 9   release_day       1125 non-null   object  
 10  release_month    1125 non-null   object  
 11  release_year     1125 non-null   object  
dtypes: object(12)
memory usage: 154.3+ KB
```

In [122]: `movie_collection = pd.concat([common_tconst,outer_sect])`

executed in 14ms, finished 13:41:12 2021-05-01

In [123]: `movie_collection.head()`

executed in 15ms, finished 13:41:12 2021-05-01

Out[123]:

	tconst	Title	Genres	runtime_minutes	vote_average	vote_count	budget_est	wor
0	tt0892769	How to Train Your Dragon	[Fantasy, Adventure, Animation, Family]	98	8.09508	618909	1.65e+08	
1	tt1228705	Iron Man 2	[Adventure, Action, Science Fiction]	124	6.99631	670058	1.85e+08	
2	tt1375666	Inception	[Action, Science Fiction, Adventure]	148	8.79405	1.86325e+06	1.6e+08	
3	tt0814255	Percy Jackson & the Olympians: The Lightning T...	[Adventure, Fantasy, Family]	118	5.90499	169337	9.5e+07	
4	tt0435761	Toy Story 3	[Animation, Family, Comedy]	103	8.29275	690558	2e+08	

```
In [124]: movie_collection = movie_collection.reset_index()
movie_collection = movie_collection.drop(columns = 'index')
movie_collection.info()
executed in 15ms, finished 13:41:13 2021-05-01
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2373 entries, 0 to 2372
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   tconst          2373 non-null    object  
 1   Title            2373 non-null    object  
 2   Genres           2373 non-null    object  
 3   runtime_minutes  2373 non-null    object  
 4   vote_average     2373 non-null    object  
 5   vote_count       2373 non-null    object  
 6   budget_est       2373 non-null    object  
 7   worldwide_gross  2373 non-null    object  
 8   release_week_day 2373 non-null    object  
 9   release_day      2373 non-null    object  
 10  release_month    2373 non-null    object  
 11  release_year     2373 non-null    object  
dtypes: object(12)
memory usage: 222.6+ KB
```

```
In [125]: movie_collection.tconst.nunique()
executed in 14ms, finished 13:41:14 2021-05-01
```

Out[125]: 2373

Sick! The dataframe is basically ready. For some reason all the columns are of object datatype, let's see if we can cast the numeric columns differently.

```
In [126]: movie_collection = movie_collection.astype({'runtime_minutes': 'int32', 'vote_ave
                                'budget_est': 'float64', 'worldwide_g
                                'release_month': 'int32', 'release_ye
executed in 14ms, finished 13:41:14 2021-05-01
```

In [127]: `movie_collection.info()`

executed in 15ms, finished 13:41:15 2021-05-01

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2373 entries, 0 to 2372
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   tconst          2373 non-null    object  
 1   Title            2373 non-null    object  
 2   Genres           2373 non-null    object  
 3   runtime_minutes  2373 non-null    int32  
 4   vote_average     2373 non-null    float64 
 5   vote_count       2373 non-null    int32  
 6   budget_est       2373 non-null    float64 
 7   worldwide_gross  2373 non-null    float64 
 8   release_week_day 2373 non-null    object  
 9   release_day      2373 non-null    int32  
 10  release_month    2373 non-null    int32  
 11  release_year     2373 non-null    int32  
dtypes: float64(3), int32(5), object(4)
memory usage: 176.2+ KB
```

In [128]: `movie_collection.to_csv('./dataFiles/movie_collection.csv', index=False)`

executed in 44ms, finished 13:41:16 2021-05-01

Woohoo! It's done. Now, in the analytics notebook we will still want to add columns such as ROI and maybe release day of the week. Additionally, we should be able to connect to different cast and crew members using tconst and joining title_principals and name_basics.